Homework 3

Problem 1

(a) I expect "Peak Brightness Date/Time (UT)" to be numerical and discrete since it is inconvenient to have times between seconds or minutes. This should be an int. I expect "Latitude (deg.)", "Longitude (deg.)" to both be numerical and continuous since there is a continnuous measurement of geographic latitude and longitude, so it should bea float. Likewise, I expect "Altitude (km)", "Velocity Components (km/s)", "Total Radiated Energy (J)","Calculated Total Impact Energy (kt)" to be numerical and continuous float data types since they are all measurements and not in seperate categories.

```
# Examine data storage types of columns in fireball.csv
import pandas as pd
df = pd.read_csv('fireball.csv')
df.dtypes
```

```
Peak Brightness Date/Time (UT)             str
Latitude (deg.)                            str
Longitude (deg.)                           str
Altitude (km)                          float64
Velocity Components (km/s)                 str
Unnamed: 5                                 str
Unnamed: 6                                 str
Total Radiated Energy (J)              float64
Calculated Total Impact Energy (kt)    float64
dtype: object
```

My predictions were, for the most part, wrong. "Peak Brightness Date/Time (UT)", "Latitude (deg.)", and "Longitude (deg.)" were all strings. Also, there are two unnamed categories that I did not see (these are the respective v_y and v_z velocity components). I also predicted "Velocity Components (km/s)" wrong as I anticipated it being a float, but it is in string format. Additionally, "Latitude (deg.)" and "Longitude (deg.)" are both partly categorical as well, categorized into the degree and also N, E, S, or W.

(b)

```
# Convert data to suitable types

# Convert the date column to datetime
df['Peak Brightness Date/Time (UT)'] = pd.to_datetime(df['Peak Brightness Date/Time (UT)'], e

# Convert the latitude and longitude columns to numeric
lat = df["Latitude (deg.)"]
df["Latitude (deg.)"] = lat.str[:-1].astype(float) * lat.str[-1].map({"N": 1, "S": -1})

lon = df["Longitude (deg.)"]
df["Longitude (deg.)"] = lon.str[:-1].astype(float) * lon.str[-1].map({"E": 1, "W": -1})

# Convert the velocity columns to numeric using their actual names
df['Velocity Components (km/s)'] = pd.to_numeric(df['Velocity Components (km/s)'], errors='co
df['Unnamed: 5'] = pd.to_numeric(df['Unnamed: 5'], errors='coerce')
df['Unnamed: 6'] = pd.to_numeric(df['Unnamed: 6'], errors='coerce')

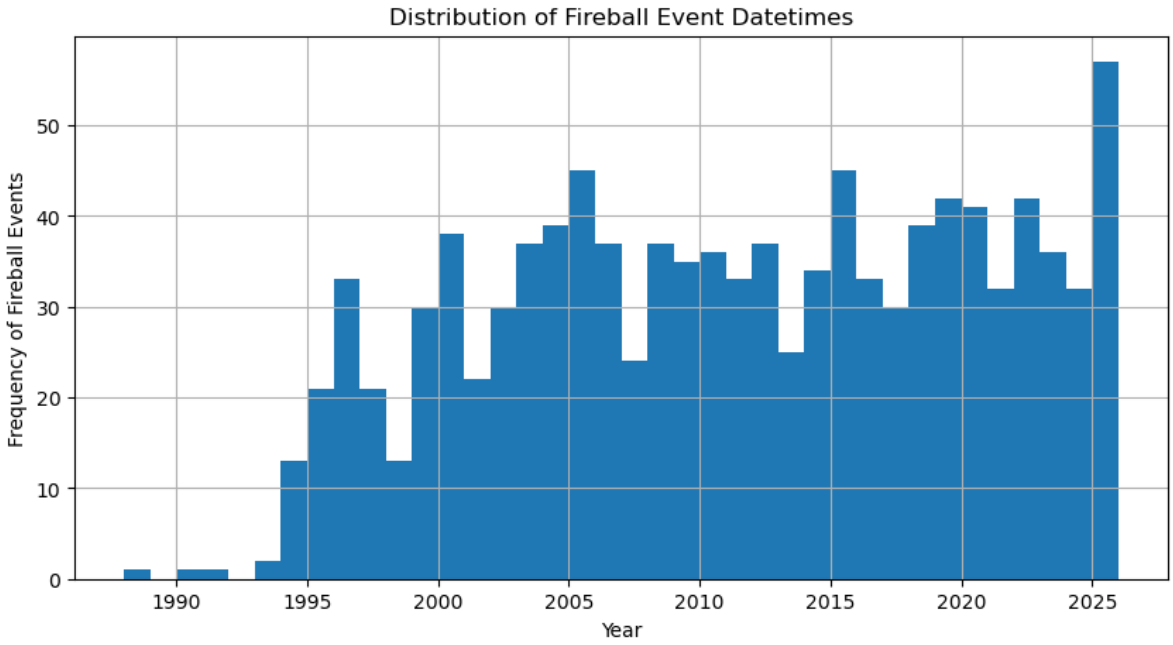
# Rename the columns
df = df.rename(columns={
    "Peak Brightness Date/Time (UT)": "datetime",
    'Velocity Components (km/s)': 'vx',
    'Unnamed: 5': 'vy',
    'Unnamed: 6': 'vz',
    "Latitude (deg.)": "lat",
    "Longitude (deg.)": "lon",
    "Altitude (km)": "altitude",
    "Total Radiated Energy (J)": "rad_energy",
    "Calculated Total Impact Energy (kt)": "impact_energy"
})

df.to_csv('fireball_cleaned.csv', index=False)
```

(c)

```
# Distribution of fireball events by year
plt = df['datetime'].dt.year.hist(bins=len(df['datetime'].dt.year.unique()), figsize=(10, 5))
plt.set_title('Distribution of Fireball Event Datetimes')
plt.set_xlabel('Year')
plt.set_ylabel('Frequency of Fireball Events')
```

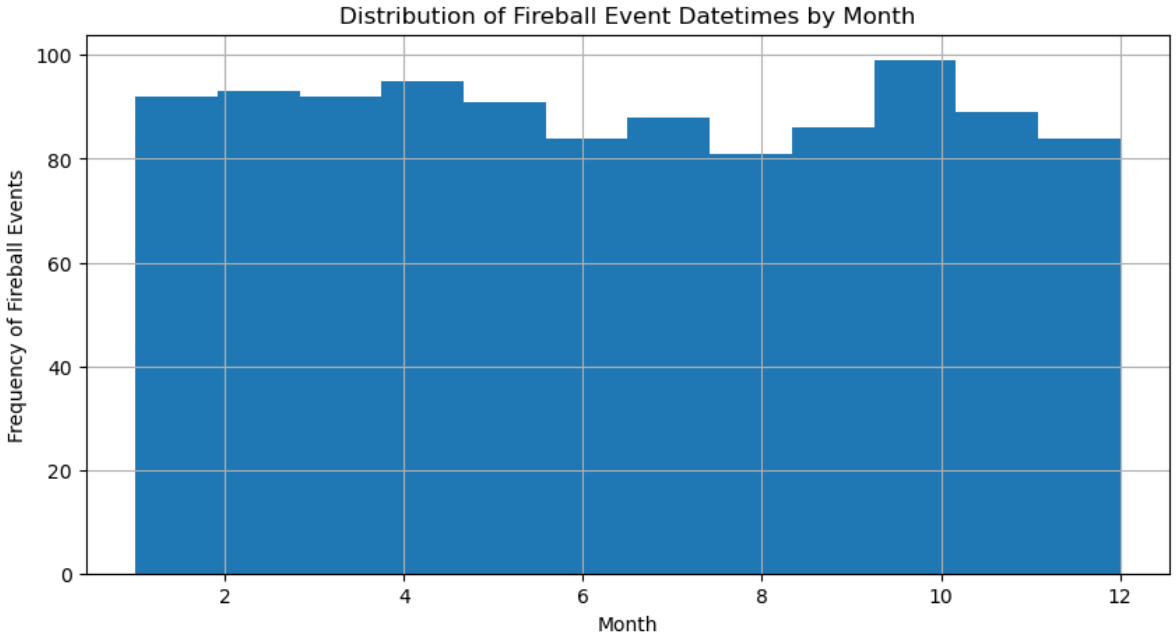
```
Text(0, 0.5, 'Frequency of Fireball Events')
```



It seems that the database did not have much data before 2000. After 2000, there seems to be a steady frequency of around 30 fireball events per year, though there do seem to be dips every 5 years or so at 2001-2002, 2007-2008, and 2013-2014.

```
# Distribution of fireball events by month
plt = df['datetime'].dt.month.hist(bins=12, figsize=(10, 5))
plt.set_title('Distribution of Fireball Event Datetimes by Month')
plt.set_xlabel('Month')
plt.set_ylabel('Frequency of Fireball Events')
```

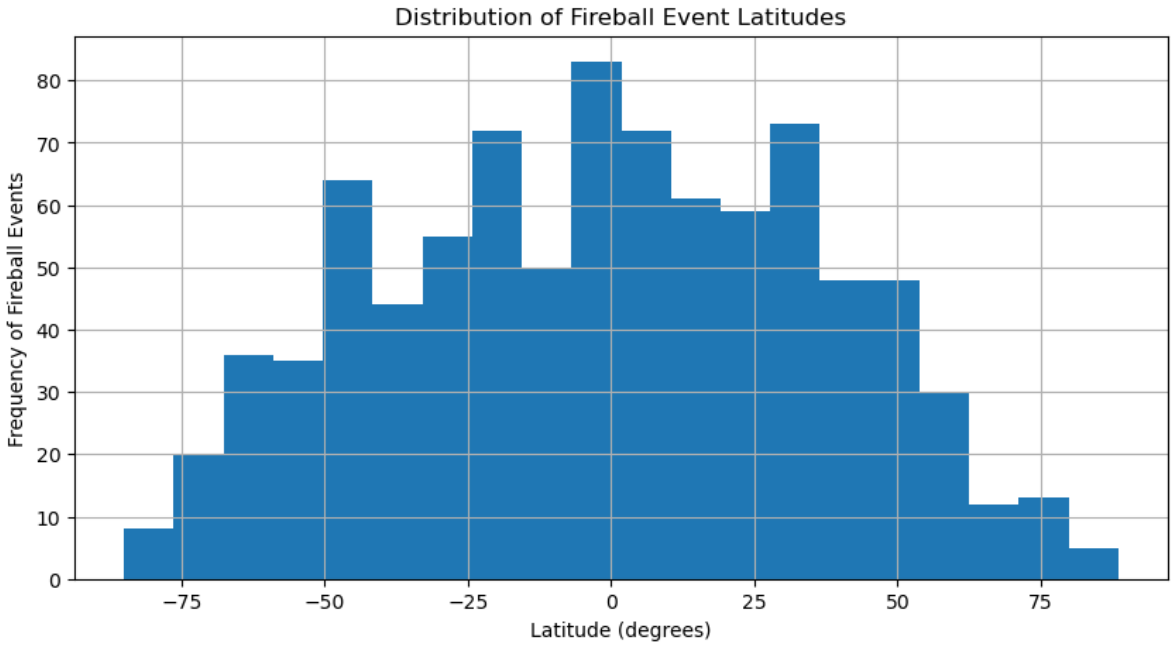
```
Text(0, 0.5, 'Frequency of Fireball Events')
```



The fireball events seem to be evenly distributed across the months, there there is a slight around late summer in August.

```
# Distribution of fireball events by latitude
plt = df['lat'].hist(bins=20, figsize=(10, 5))
plt.set_title('Distribution of Fireball Event Latitudes')
plt.set_xlabel('Latitude (degrees)')
plt.set_ylabel('Frequency of Fireball Events')
```

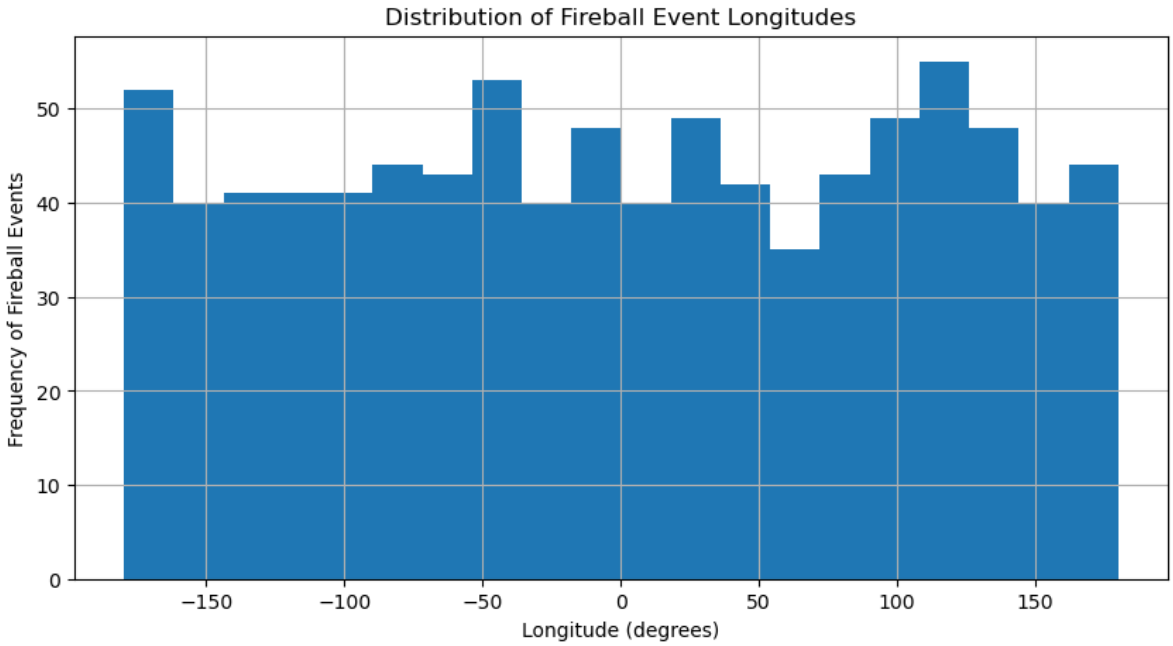
```
Text(0, 0.5, 'Frequency of Fireball Events')
```



There is a clear unimodal peak at a latitude of 0 degrees (the most impacts happen here). The further from 0 degrees we get, the less fireball events there are. This is a symmetrical distribution.

```
# Distribution of fireball events by longitude
plt = df['lon'].hist(bins=20, figsize=(10, 5))
plt.set_title('Distribution of Fireball Event Longitudes')
plt.set_xlabel('Longitude (degrees)')
plt.set_ylabel('Frequency of Fireball Events')
```

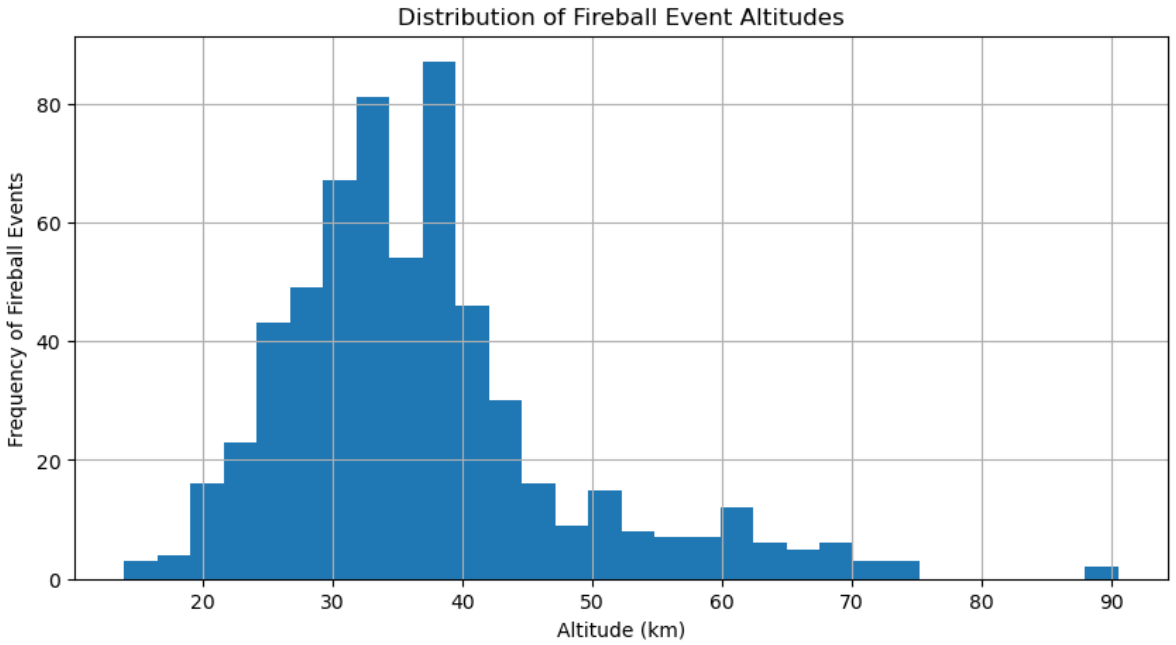
```
Text(0, 0.5, 'Frequency of Fireball Events')
```



The number of fireball events seems to be pretty uniformly destributed across longitude.

```
# Distribution of fireball events by altitude
plt = df['altitude'].hist(bins=30, figsize=(10, 5))
plt.set_title('Distribution of Fireball Event Altitudes')
plt.set_xlabel('Altitude (km)')
plt.set_ylabel('Frequency of Fireball Events')
```

```
Text(0, 0.5, 'Frequency of Fireball Events')
```



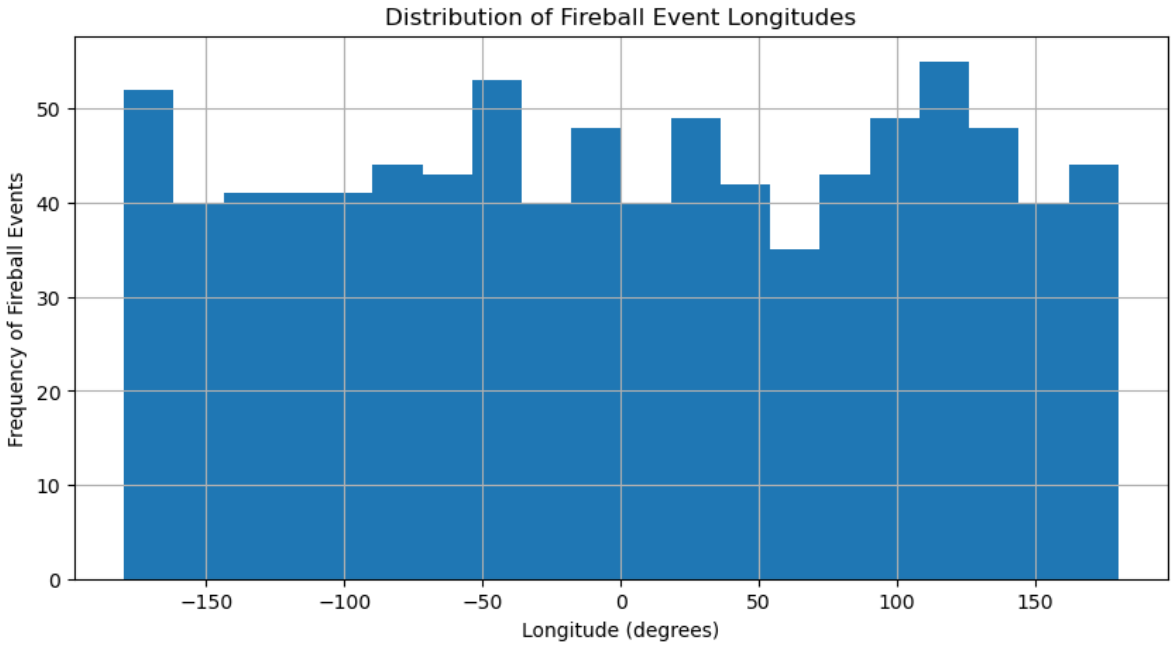
The distribution of fireball events by altitude seems to be right-skewed, withe a bimodal distribution (peaks around 30km and 40km). There is a gap around 80 km.

(d)

```
# Distribution of fireball events by longitude
plt = df['lon'].hist(bins=20, figsize=(10, 5))
plt.set_title('Distribution of Fireball Event Longitudes')
plt.set_xlabel('Longitude (degrees)')
plt.set_ylabel('Frequency of Fireball Events')

# Check for outliers in the longitude
df = pd.read_csv('fireball_cleaned.csv')
Q1 = df['lon'].quantile(0.25)
Q3 = df['lon'].quantile(0.75)
IQR = Q3 - Q1
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR
print(f"Outliers: {df[(df['lon'] < lower_bound) | (df['lon'] > upper_bound)]}")
```

```
Outliers: Empty DataFrame
Columns: [datetime, lat, lon, altitude, vx, vy, vz, rad_energy, impact_energy]
Index: []
```



The longitude is centered, at 0 degrees with a pretty even spread. There seem to be no outliers based purely on the shape of the histogram shape, and using the 1.5*IQR rule, we see that there are no outliers in the longitude. This distribution is consistent with what expect. Moreover, since fireball events orbit the earth, I expect all longitudes (measured vertically) to be evenly covered. It makes sense that the histogram mirrors a uniform distribution.

(e) Make a Q-Q plot or an equivalent plot of the distribution of the latitudes of fireballs relative to the theoretical distribution that fireballs occur at a density proportional to the cosine of latitude.

```
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats
from statsmodels.graphics.gofplots import qqplot

# Signed coordinates: N and E positive, S and W negative
d = df.dropna(subset=['lat', 'lon'])
lat = d['lat']
lon = d['lon']

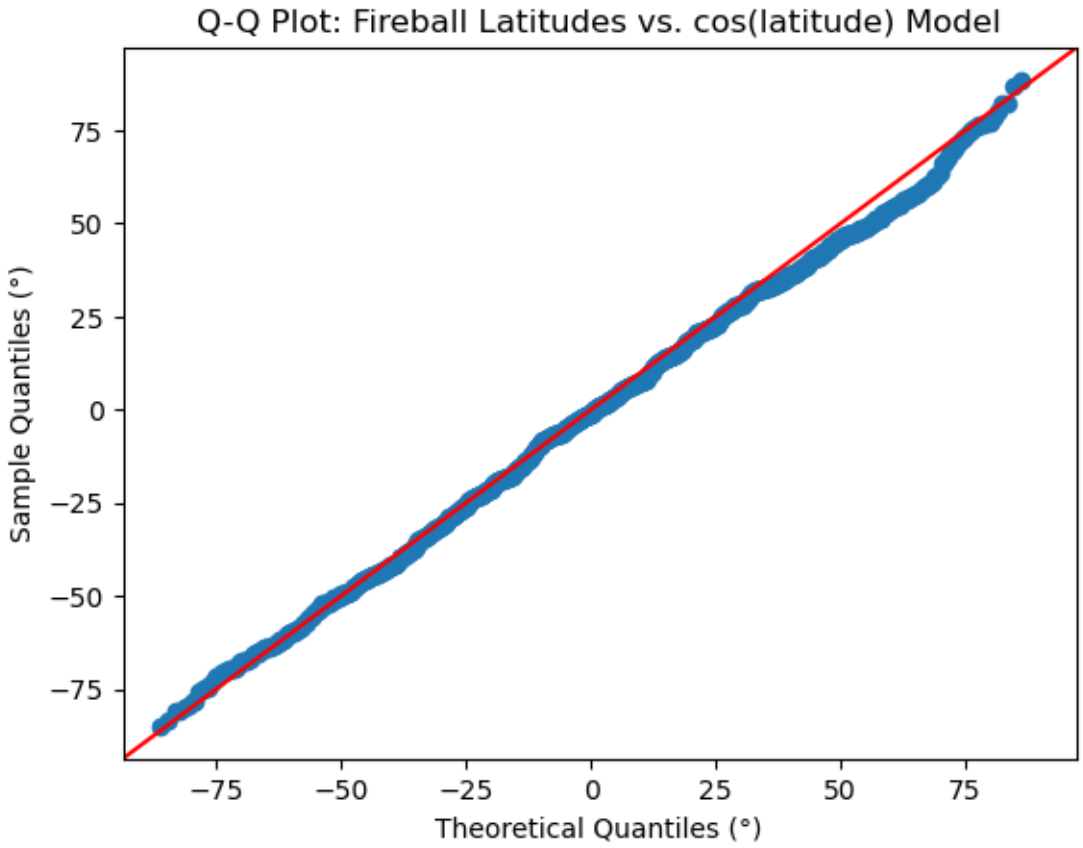
# Theoretical distribution: density proportional to cos(latitude), in degrees
class coslat_gen(stats.rv_continuous):
    def _pdf(self, x): return np.cos(np.radians(x)) * np.pi / 360
    def _cdf(self, x): return (1 + np.sin(np.radians(x))) / 2
```

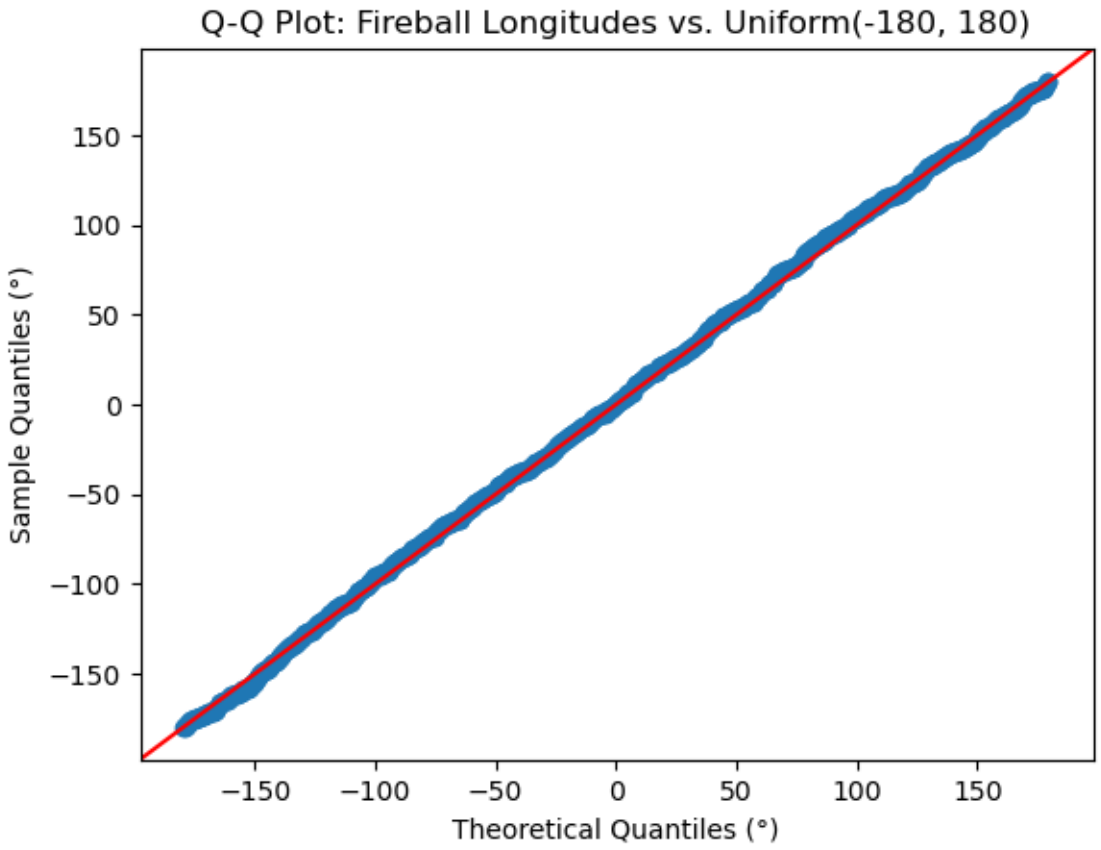
```
    def _ppf(self, p): return np.degrees(np.arcsin(2 * p - 1))

coslat = coslat_gen(a=-90, b=90, name='coslat')

# Q-Q plot for latitude
qqplot(lat, dist=coslat, line='45')
plt.title('Q-Q Plot: Fireball Latitudes vs. cos(latitude) Model')
plt.xlabel('Theoretical Quantiles (°)')
plt.ylabel('Sample Quantiles (°)')
plt.show()

# Q-Q plot for longitude
qqplot(lon, dist=stats.uniform, loc=-180, scale=360, line='45')
plt.title('Q-Q Plot: Fireball Longitudes vs. Uniform(-180, 180)')
plt.xlabel('Theoretical Quantiles (°)')
plt.ylabel('Sample Quantiles (°)')
plt.show()
```





(f)

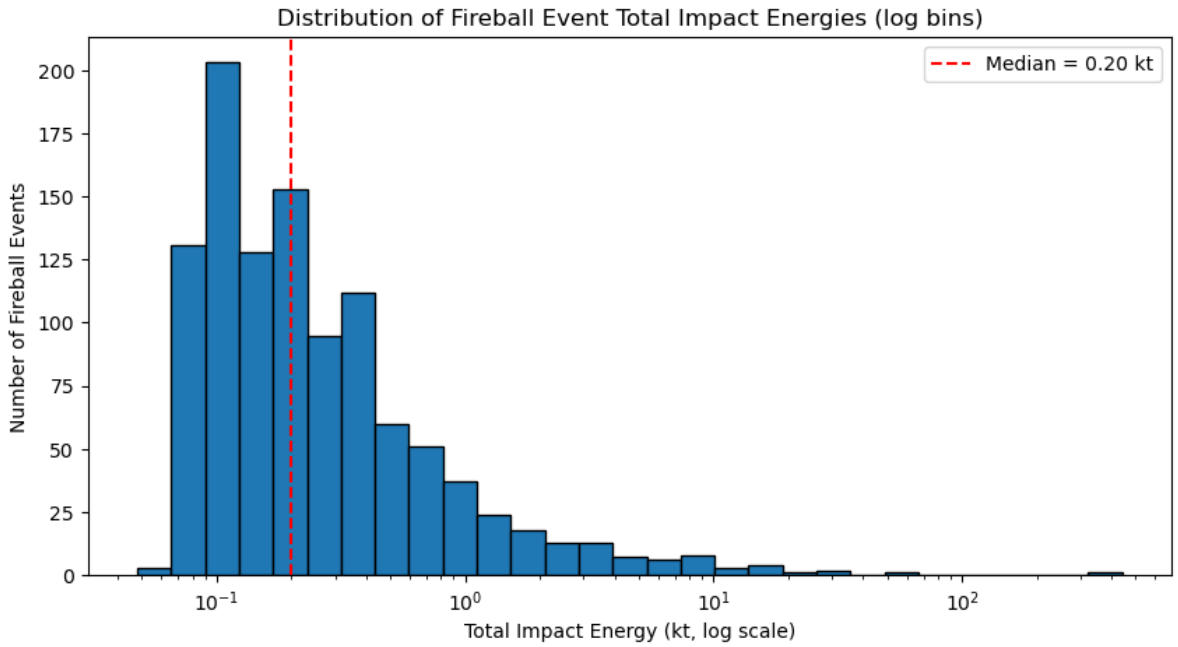
```
import numpy as np
import matplotlib.pyplot as plt

e = df['impact_energy'].dropna()
e = e[e > 0]  # log scale requires positive values

# 30 bins evenly spaced in log10(energy)
bins = np.logspace(np.log10(e.min()), np.log10(e.max()), 30)

plt.figure(figsize=(10, 5))
plt.hist(e, bins=bins, edgecolor='black')
plt.xscale('log')
plt.axvline(e.median(), color='red', linestyle='--', label=f'Median = {e.median():.2f} kt')
plt.title('Distribution of Fireball Event Total Impact Energies (log bins)')
plt.xlabel('Total Impact Energy (kt, log scale)')
plt.ylabel('Number of Fireball Events')
```

```
plt.legend()
plt.show()
```




```
q1, med, q3 = e.quantile([0.25, 0.5, 0.75])
print(f"Median: {med:.3f} kt, Q1: {q1:.3f}, Q3: {q3:.3f}, IQR: {q3 - q1:.3f}")
print(f"Mean: {e.mean():.3f} kt (pulled up by the right tail)")

# Outliers using the 1.5*IQR rule on log10(energy)
le = np.log10(e)
lq1, lq3 = le.quantile([0.25, 0.75])
liqr = lq3 - lq1
outliers = e[(le < lq1 - 1.5 * liqr) | (le > lq3 + 1.5 * liqr)]
print(outliers.sort_values())
```

```
Median: 0.200 kt, Q1: 0.110, Q3: 0.410, IQR: 0.300
Mean: 1.124 kt (pulled up by the right tail)
910       3.1
784       3.2
857       3.3
59        3.4
627       3.5
467       3.5
```

```
574       3.8
103       3.8
157       3.8
793       3.9
399       3.9
279       4.3
819       4.6
543       4.9
1059      5.0
718       5.0
73        5.1
1073      5.2
263       6.0
112       6.5
487       6.9
162       7.0
108       7.1
935       7.3
446       7.6
325       7.9
865       8.8
864       8.8
895       9.0
559       9.7
965       9.8
481       9.8
202      12.0
382      13.0
783      13.0
691      14.0
579      14.0
1074     14.0
780      18.0
598      20.0
1069     30.0
605      33.0
284      49.0
485     441.0
Name: impact_energy, dtype: float64
```

The distribution of impact energies of the fireball events appears still right skewed, with a unimodal peak around 10^{-1}. It is centered at around a median of 0.200 kilotons, with a

spread of 0.300 kilotons (the IQR). By the 1.5*IQR outlier rule, there are quite a few, with the event above 400 kilotons as the most obvious outlier. There are also several events with below 10 kt of energy.

(g)

```
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats

d = df.dropna(subset=['lon'])
lon = d['lon']

# Roundign
deg = np.floor(lon).astype(int)
deg[deg == 180] = -180

# Count fireballs
counts = np.bincount(deg + 180, minlength=360)   # length 360

# Histogram-of-the-histogram
k = np.arange(counts.max() + 1)
observed = np.bincount(counts, minlength=len(k))

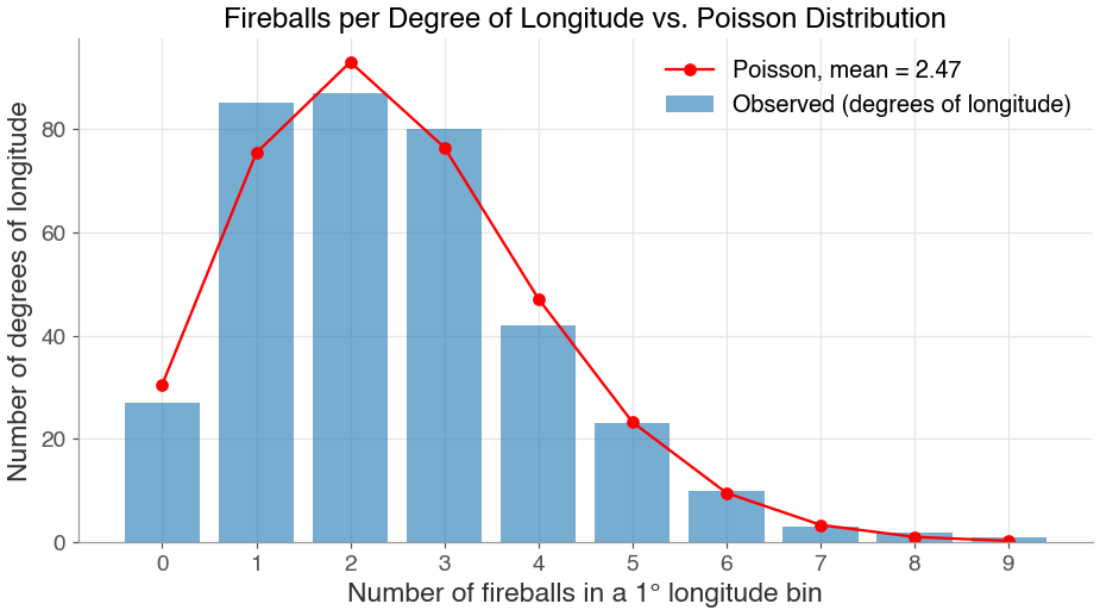
mu = counts.mean()           # should be about 888/360 = 2.47
expected = 360 * stats.poisson.pmf(k, mu)

print("Degrees plotted:", observed.sum())   # should be 360
print("Mean events per degree:", mu)
print("Variance:", counts.var(), "(Poisson variance = mean)")
print("Empty degrees:", observed[0], "expected:", round(expected[0], 1))

plt.figure(figsize=(10, 5))
plt.bar(k, observed, width=0.8, alpha=0.6, label='Observed (degrees of longitude)')
plt.plot(k, expected, 'ro-', label=f'Poisson, mean = {mu:.2f}')
plt.xticks(k)
plt.xlabel('Number of fireballs in a 1° longitude bin')
plt.ylabel('Number of degrees of longitude')
plt.title('Fireballs per Degree of Longitude vs. Poisson Distribution')
plt.legend()
plt.show()
```

```
Degrees plotted: 360
```

```
Mean events per degree: 2.466666666666667
Variance: 2.5711111111111107 (Poisson variance = mean)
Empty degrees: 27 expected: 30.6
```



The poisson distribution seems to agree pretty well with the number of fireballs in each degree of longitude. There are slightly lower amount of number of degrees of longitude compared to the discrete poisson distribution's number of degrees of longitude for 0, 2, and 4 fireball events and a slightly higher amount of degrees of longitude compared to the amount of degrees of longitude which the poisson distribution predicts for and 3 fireball events in each 1 degree longitude bin. Otherwise, the poisson ditribution seems to match the observed events pretty well.

Problem 2

Visualization 1 & 2: Do more Heating Degree Days correlate to more snowfall on a year-to-year basis?

That is, do colder ("freezing") days mean it is more likely to see snow? Does it mean there is deeper snowfall measured?

```
# Visualization 1: Snowfall each winter over time, colored by freezing days
anom = seasons['FREEZE'] - mean_f
lim = np.abs(anom).max()
norm = TwoSlopeNorm(vcenter=0, vmin=-lim, vmax=lim)
cmap = plt.cm.coolwarm_r   # blue = more freezing days, red = fewer

fig, ax = plt.subplots(figsize=(13, 9))
fig.subplots_adjust(left=0.07, right=0.98, top=0.80, bottom=0.30)

ax.bar(seasons.index, seasons['SNOW'], color=cmap(norm(anom)),
       edgecolor='white', linewidth=0.8, width=0.85)
ax.axhline(mean_s, color=GRAY, linestyle='--', linewidth=1.4,
           label=f"{n}-season average: {mean_s:.1f} in")
ax.plot(seasons.index, snow_tr.intercept + snow_tr.slope * seasons.index,
        color=DARK, linestyle=':', linewidth=1.8,
        label=f"Trend: {snow_tr.slope*10:+.1f} in/decade")

ax.set_xticks(ticks)
ax.set_xticklabels([label(s) for s in ticks], rotation=45, ha='right')
ax.set_xlabel('Winter season')
ax.set_ylabel('Total snowfall (inches)')
plt.show()

# Visualization 2: Freezing days each winter over time
x, y = seasons.index.values, seasons['FREEZE'].values

fig, ax = plt.subplots(figsize=(13, 8.5))
fig.subplots_adjust(left=0.07, right=0.98, top=0.80, bottom=0.30)

# Shade above/below the average
ax.fill_between(x, y, mean_f, where=y >= mean_f, interpolate=True, color=BLUE, alpha=0.18)
ax.fill_between(x, y, mean_f, where=y < mean_f, interpolate=True, color=RED, alpha=0.18)
ax.plot(x, y, color=DARK, linewidth=1.8, zorder=2)

ax.axhline(mean_f, color=GRAY, linestyle='--', linewidth=1.4)
ax.plot(x, freeze_tr.intercept + freeze_tr.slope * x, color=DARK,
        linestyle=':', linewidth=1.8)

# Label the coldest and mildest winters
coldest, mildest = seasons['FREEZE'].idxmax(), seasons['FREEZE'].idxmin()
for s, offset in [(coldest, 12), (mildest, -26)]:
    ax.annotate(f"{label(s)}: {seasons.loc[s, 'FREEZE']} days",
```

```
                (s, seasons.loc[s, 'FREEZE']), xytext=(0, offset),
                textcoords='offset points', ha='center', fontsize=NOTE_SIZE, color=DARK)

ax.set_xticks(ticks)
ax.set_xticklabels([label(s) for s in ticks], rotation=45, ha='right')
ax.set_xlabel('Winter season')
ax.set_ylabel('Freezing days (high  32°F)')
ax.set_ylim(0, y.max() * 1.45)
```

```
Dropped seasons (incomplete data): [1997, 1998]
Number of seasons: 28
```

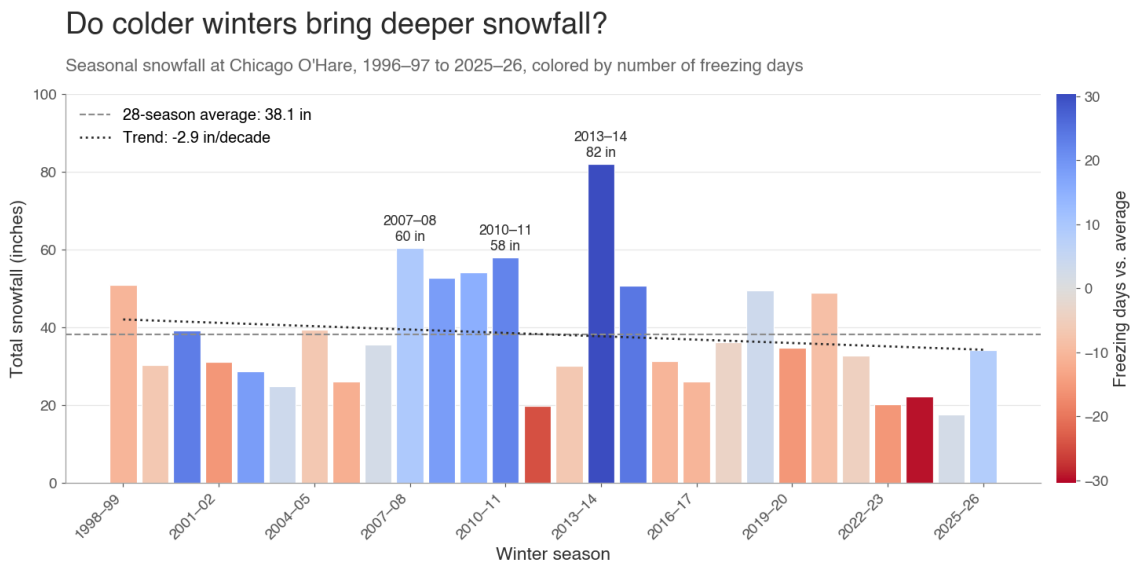


Figure 1. Each bar is the total snowfall at Chicago O'Hare for one winter season. Bar color shows how many freezing days (daily high ≤ 32°F) that winter had compared with the 28-season average of 42 days: blue winters had more freezing days, red winters had fewer. Snowfall averaged 38.1 inches per season. Across all seasons, freezing days and snowfall have a correlation of r = 0.66 (p = 0.000). The 1996–97 and 1997–98 seasons were excluded because there was not enough data. Data: NOAA GHCN-Daily, Chicago O'Hare International Airport.

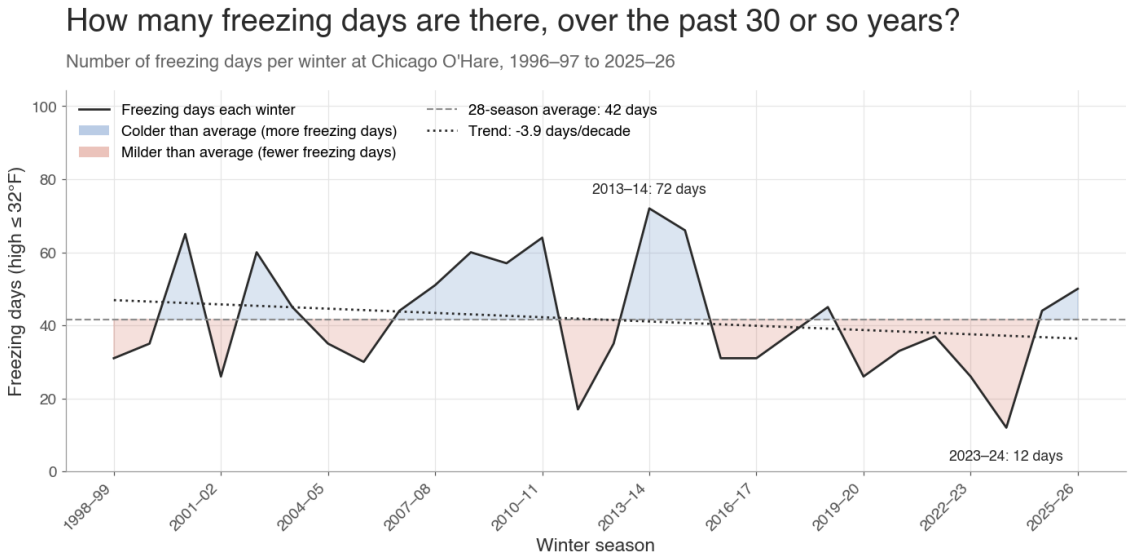


Figure 2. Number of days in each winter season when the high temperature stayed at or below 32°F. Shading marks winters above (blue) and below (red) the 28-season average of 42 days. The dotted line is a linear trend of -3.9 days per decade. The 1996–97 and 1997–98 seasons were excluded because there was not enough data. The coldest winter, 2013–14, had 72 freezing days; the mildest, 2023–24, had 12. Data: NOAA GHCN-Daily, Chicago O'Hare International Airport.

Visualization 3 (hourly): When does it snow? Snow by time of day

```
frames = []
for yr in range(2020, 2026):
    start = "2020-07-01" if yr == 2020 else f"{yr}-01-01"
    end   = "2025-06-30" if yr == 2025 else f"{yr}-12-31"
    url = ("https://www.ncei.noaa.gov/access/services/data/v1"
           "?dataset=local-climatological-data&stations=72530094846"
           f"&startDate={start}&endDate={end}&format=csv")
    frames.append(pd.read_csv(url, low_memory=False))
    print(yr, "downloaded")
raw = pd.concat(frames, ignore_index=True)

# Hour by hour reports
hr = raw[raw['REPORT_TYPE'].astype(str).str.strip() == 'FM-15'].copy()
hr['DATE'] = pd.to_datetime(hr['DATE'])
hr = hr[hr['DATE'].dt.month.isin([11, 12, 1, 2, 3])]
hr['hour'] = hr['DATE'].dt.hour
hr['snowing'] = hr['HourlyPresentWeatherType'].astype(str).str.contains(r'(?<!BL)SN', regex=
```

```
# % of hours with snow at each hour of the day
by_hour = hr.groupby('hour')['snowing'].mean() * 100
peak_h, low_h = by_hour.idxmax(), by_hour.idxmin()
hlabel = lambda h: f"{h % 12 or 12} {'AM' if h < 12 else 'PM'}"
fig, ax = plt.subplots(figsize=(13, 8.5))
fig.subplots_adjust(left=0.08, right=0.98, top=0.80, bottom=0.28)

# Background bands for parts of the day
bands = [(-0.5, 5.5, 'Night'), (5.5, 11.5, 'Morning'),
         (11.5, 17.5, 'Afternoon'), (17.5, 23.5, 'Evening')]
ymax = by_hour.max() * 1.3
```

```
2020 downloaded
2021 downloaded
2022 downloaded
2023 downloaded
2024 downloaded
2025 downloaded
```

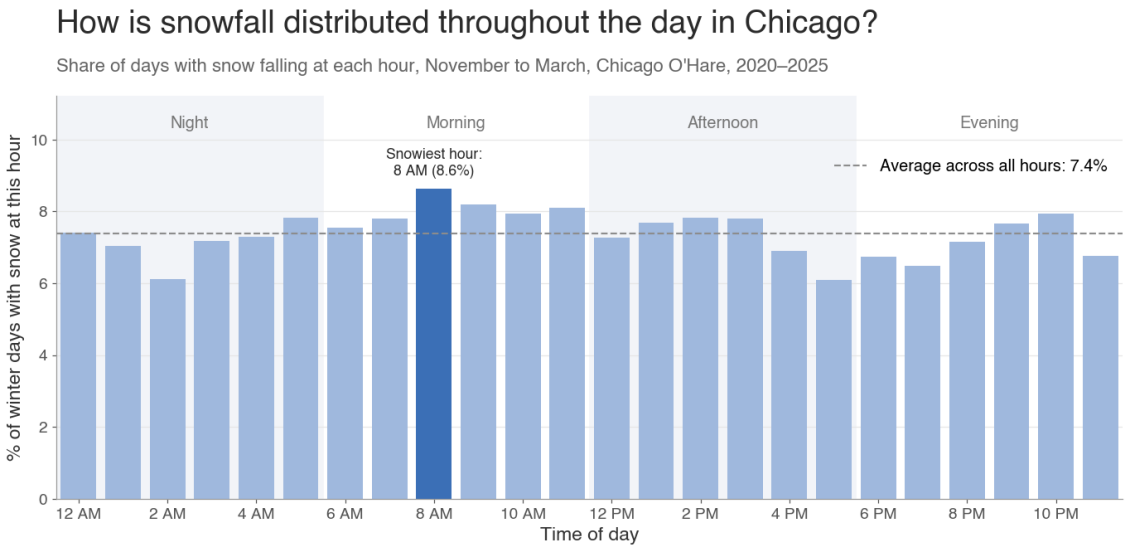


Figure 3. Each bar shows the percentage of November–March days on which snow was falling at that hour, based on hourly weather observations at Chicago O'Hare from July 2020 to June 2025. For example, a bar at 8% means snow was falling at that hour on about 8% of winter days. Across all hours, the average was 7.4%. Snow was most common around 8 AM (8.6% of days) and least common around 5 PM (6.1%).Data: NOAA Local Climatological Data, Chicago O'Hare.

```
hour
0     7.4
1     7.0
2     6.1
3     7.2
4     7.3
5     7.8
6     7.5
7     7.8
8     8.6
9     8.2
10    7.9
11    8.1
12    7.3
13    7.7
14    7.8
15    7.8
16    6.9
17    6.1
18    6.7
19    6.5
20    7.2
21    7.7
22    7.9
23    6.8
Name: snowing, dtype: float64
```

AI Statement

For Question 1 (fireballs), Claude helped me write and debug code for the Q-Q plots of latitude and longitude, the log-binned histogram of total impact energies, and the comparison of fireballs per degree of longitude with a Poisson distribution, and it helped me interpret the results. For Question 2 (Chicago weather), I chose the relationships I wanted to analyze about freezing days/HDD. Claude then helped me access NOAA data for Chicago O'Hare and generated first drafts of the code for the visualizations. I revised the plots, with Claude's help on some changes, by rewording the titles and axis labels, enlarging the fonts, adjusting the spacing, simplifying the freezing-days plot (removing the data points and adding a key), and adding notes to the captions about seasons excluded for missing data. I take full responsibility for the accuracy of the submitted work.